



**POLITECNICO**  
MILANO 1863

# **Actor Model - Akka**

Luca Mottola

Credits: Alessandro Margara

luca.mottola@polimi.it

<http://mottola.faculty.polimi.it>

# Outline

---

- Fundamentals
- Communication
- Fault-tolerance

# Fault-tolerance

# Managing faults


---

- Faults and exceptional behaviors are managed through the concept/pattern of **supervision**
- Each actor has a supervisor that monitors its execution state
  - ...which is plainly another actor
- A supervisor can decide to terminate and possibly restart a supervised faulty actor
  - Restart from scratch or from the latest available state

# Supervisor hierarchy

---


- Applications are typically organized in **supervision trees**
- Each supervisor knows how to handle failures in its directly supervised nodes
- If a supervisor cannot handle a problem locally, it propagates the fault to the upper layer
- On the top of the hierarchy there are typically standard supervisors offered by the actor framework



# Supervisor hierarchy (again)

---

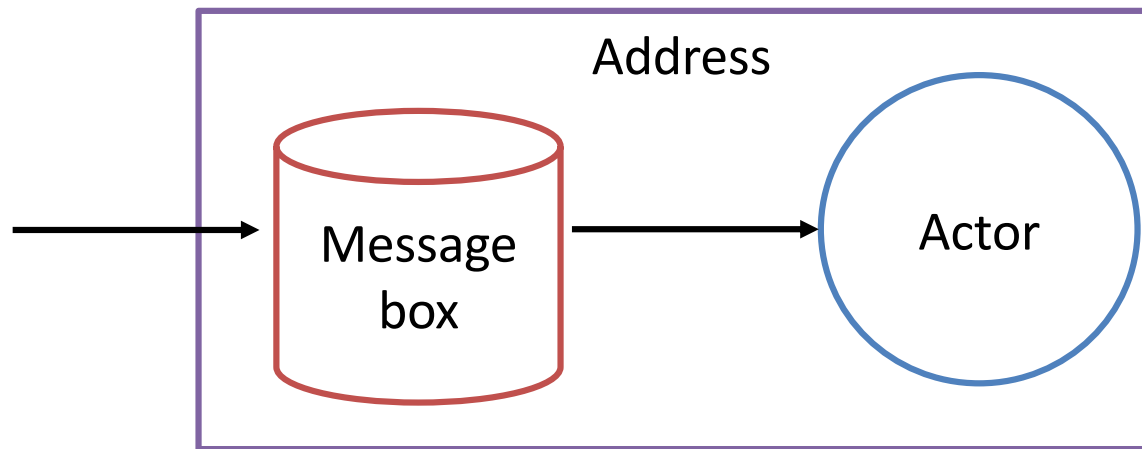
- The model favors the definition of hierarchies of responsibility
- If a **failing actor** contains important data that shall not be lost, the **supervisor actor** sources any possibly dangerous subtask to children



# Supervisor

---

- The framework handles the lifecycle of an actor
- When an actor is restarted
  - Its address does not change
  - Its message box is retrieved from persistent state
    - Which lives outside of the actor memory



Hands-on: Akka



# Fault tolerance

---


- An actor is responsible for all actors it creates in its context
  - Created through the `getContext().actorOf()` method
  - Stopped through the `getContext().stop()` method
- The supervision strategy can be customized by overriding the **`supervisorStrategy()`** method
  - Depending on the type of exception
  - Depending on the number of errors
  - Independent of the supervised child
    - Philosophy: if you need more flexibility, build a deeper hierarchy!
    - An actor should normally supervise actors of the same type, that is, providing the same or similar functionality

how we do supervision need to be independent to whom we

# Fault tolerance

---

- Different supervision strategies exist
  - One-for-one
  - One-for-all
- Similarly, different actions may be applied when a fault occurs, and the supervisor is asked to manage the situation
  - **stop()**, **restart()**, **resume()**, **escalate()**



# Example

---

```
public class CounterSupervisorActor extends AbstractActor {
```

```
    private static SupervisorStrategy strategy =
```

```
        new OneForOneStrategy(
```

```
            10, 10 = maximum number of faults that may be managed with this strategy within a
```

```
            Duration.ofMinutes(1),
```

```
            DeciderBuilder.match(Exception.class, e ->  
                SupervisorStrategy.restart())
```

```
                .build());
```

what we need to do upon recognising a fault is specified with the match() clause: that says whenever a g

```
@Override
```

```
public SupervisorStrategy supervisorStrategy() {
```

```
    return strategy;
```

```
...
```

# Example

---

```
public class CounterActor extends AbstractActor {

...

    void onMessage(DataMessage msg) throws Exception {
        if (msg.getCode() == Counter.NORMAL_OP) {
            System.out.println("I am executing a NORMAL operation...counter is now " + (++counter));
        } else if (msg.getCode() == Counter.FAULT_OP) {
            System.out.println("I am emulating a FAULT!");
            throw new Exception("Actor fault!");
        }
    }

    @Override
    public void preRestart(Throwable reason, Optional<Object> message) {
        System.out.print("Preparing to restart...");
    }

    @Override
    public void postRestart(Throwable reason) {
        System.out.println("...now restarted!");
    }

    static Props props() {
        return Props.create(CounterActor.class);
    }
}
```

# Example

---

```
public class CounterSupervisor {

    public static final int NORMAL_OP = 0;
    public static final int FAULT_OP = -1;

    public static final int FAULTS = 1;

    public static void main(String[] args) {
        scala.concurrent.duration.Duration timeout = scala.concurrent.duration.
                                                    Duration.create(5, SECONDS);

        final ActorSystem sys = ActorSystem.create("System");
        final ActorRef supervisor = sys.actorOf(CounterSupervisorActor.props(), "supervisor");

        ActorRef counter;
        try {
            scala.concurrent.Future<Object> waitingForCounter = ask(supervisor,
                                                                    Props.create(CounterActor.class), 5000);
            counter = (ActorRef) waitingForCounter.result(timeout, null);

            counter.tell(new DataMessage(NORMAL_OP), ActorRef.noSender());
            for (int i = 0; i < FAULTS; i++)
                counter.tell(new DataMessage(FAULT_OP), ActorRef.noSender());
            counter.tell(new DataMessage(NORMAL_OP), ActorRef.noSender());

            sys.terminate();
        }
    }
}
```

...

# Clustering

---

- Akka clustering offers a membership service
  - Decentralized
  - No single point of failure/bottleneck
- Implementation
  - Peer to peer
  - Gossip protocol
  - Automatic failure detection

# Clustering terminology

---

- **Node:** a logical member of a cluster
  - There can be multiple nodes on each physical machine
  - Each node is identified by a tuple hostname:port:uid
  - You can think of each node as a process / actor system
- **Cluster:** a set of nodes joined together through the membership service
- **Leader:** a role in the cluster
  - A single node in the cluster acts as a leader
  - The leader manages cluster convergence

# Clustering basics

---

- An Akka application can be distributed over a cluster
- Each node hosts some part of the application
- Cluster membership and the actors running on a member node are decoupled
- A node can be a member of a cluster hosting any number of actors
- An actor system/node joins a cluster by sending a **join command** to one of the nodes in the cluster



# Clustering protocol

---

- Nodes organize themselves into an overlay network
- They distribute information about cluster members using a gossip protocol
  - Nodes propagate messages containing their current view of the membership
  - Nodes update their view based on the received messages
    - Messages are designed in such a way that the state of nodes eventually converges
    - They contain **vector clocks** to record a partial order of the events (nodes joining, leaving, ...) observed in the environment

# Clustering protocol

---

- Information about the cluster converges at a given node when the node can prove that the cluster state it is observing **is seen the same** by all other nodes in the same cluster
- Gossip convergence cannot occur while some node is “unreachable”
  - A state in the lifecycle of nodes indicating that it is not currently possible to communicate with the node
  - The node need to become reachable again or be removed from the cluster

# Seed nodes

---

- Seed nodes are configured contact points for new nodes that want to join the cluster
- When a new node wants to join the cluster
  - It contacts all the seed nodes
    - At least one needs to be active
  - It sends a join command to the first seed node that answers

# Cluster tools

---

- Akka offers higher-level tools that build on top of clustering
  - Cluster singleton: to ensure that a single actor of a certain type exists in the cluster
  - Cluster sharding: distributes actors across nodes of the cluster
    - Ensuring that they can communicate without knowing their physical location
  - Distributed data: creates a distributed key-value store across the nodes of the cluster
  - Cluster metrics: to publish and collect health metrics about cluster members

# Further readings

---

- Akka actors official documentation
  - <https://doc.akka.io/docs/akka/current/index-actors.html>